

Solution to question 5

```
In [98]: 1 import numpy as np
2 import sklearn
3 from sklearn.datasets import fetch_20newsgroups
4 from sklearn.feature_extraction.text import TfidfVectorizer
5 from sklearn.feature_extraction.text import CountVectorizer
6 from sklearn.pipeline import Pipeline
7 from sklearn.naive_bayes import MultinomialNB
8 from sklearn.model_selection import train_test_split
9 from sklearn.cluster import KMeans
10
11 twenty_train = fetch_20newsgroups(subset='train', shuffle=True)
12 twenty_test = fetch_20newsgroups(subset='test', shuffle=True)
13
14 print(twenty_train.keys())
15
```

```
dict_keys(['data', 'filenames', 'target_names', 'target', 'DESCR'])
```

Using the guide given in <https://towardsdatascience.com/machine-learning-nlp-text-classification-using-scikit-learn-python-and-nltk-c52b92a7c73a> (<https://towardsdatascience.com/machine-learning-nlp-text-classification-using-scikit-learn-python-and-nltk-c52b92a7c73a>) The data is in key 'data' with the target labels in key 'target'

As sklearn already provides the data in the form of train and test data, using 300 samples from the training data for now

```
In [99]: 1 data = twenty_train.data
2 data = data[0:300]
```

Removed the stop words while using TfidfVectorizer

```
In [100]: 1 vectorizer = TfidfVectorizer(stop_words = 'english')
2 X = vectorizer.fit_transform(data)
3 # print(vectorizer.get_feature_names())
```

```
In [101]: 1 true_k = 5
2 model = KMeans(n_clusters=true_k, init='k-means++', max_iter=100, n
3 model.fit(X)
```

```
Out[101]: KMeans(algorithm='auto', copy_x=True, init='k-means++', max_iter=100,
n_clusters=5, n_init=1, n_jobs=None, precompute_distances='aut
random_state=None, tol=0.0001, verbose=0)
```

```
In [102]: 1 print("Top terms per cluster:")
2 order_centroids = model.cluster_centers_.argsort()[:, :-1]
3 terms = vectorizer.get_feature_names()
4 for i in range(true_k):
5     print("\nCluster %d:" % i),
6     for ind in order_centroids[i, :1]:
7         print(' %s' % terms[ind]),
8     print
```

Top terms per cluster:

Cluster 0:
edu

Cluster 1:
edu

Cluster 2:
edu

Cluster 3:
edu

Cluster 4:
com

```
In [103]: 1 print("\n")
2 print("Prediction")
3
4 Y = vectorizer.transform(["Android surpasses iOS."])
5 prediction = model.predict(Y)
6 print("the class of the above news is ", twenty_train.target_names[prediction])
7
8 Y = vectorizer.transform(["Joe Biden defeats Trump"])
9 prediction = model.predict(Y)
10 print("the class of the above news is ", twenty_train.target_names[prediction])
```

Prediction

the class of the above news is comp.graphics

the class of the above news is comp.graphics

I am not sure if it is working correctly as it classifies the political news as ms-windows.misc

Solution to question 6, based on examples provided in
<http://www.science.smith.edu/~jcrouser/SDS293/labs/lab7-py.html>
<http://www.science.smith.edu/~jcrouser/SDS293/labs/lab7-py.html>

```
In [104]: 1 import pandas as pd
2 from sklearn.model_selection import cross_val_score
3 from sklearn.model_selection import train_test_split
4 from sklearn.linear_model import LinearRegression
5 from sklearn.linear_model import LogisticRegression
6 from sklearn.metrics import confusion_matrix
7 data = pd.read_excel("ENB2012_data.xlsx")
```

```
In [105]: 1 data.head()
```

```
Out[105]:
```

	X1	X2	X3	X4	X5	X6	X7	Y
0	0.74	686.0	245.0	3.5	2	0.0	0	0
1	0.74	686.0	245.0	3.5	4	0.0	0	0
2	0.74	686.0	245.0	3.5	5	0.0	0	0
3	0.74	686.0	245.0	3.5	3	0.0	0	0
4	0.71	710.5	269.5	3.5	2	0.0	0	0

```
In [106]: 1 X = data.iloc[:,1:6]
2 y = data.iloc[:,7]
```

```
In [107]: 1 x_train, x_test, y_train, y_test = train_test_split(X, y)
```

```
In [108]: 1 clf = LinearRegression()
2 scores = cross_val_score(clf, x_train, y_train, cv=5)
3 # scores = cross_val_score(clf, X, y, cv=5)
```

```
In [109]: 1 print("average cross validation accuracy is ",scores.mean())

average cross validation accuracy is  0.9139580130418947
```

After cross validation, the classifier performs 90% correctly in the training data. Now the classifier will be tested on the testing data. The data will be fitted to the x_train and y_train and confusion matrix will be shown w.r.t the testing data.

```
In [110]: 1 clf.fit(x_train, y_train)
2 y_pred_linear = clf.predict(x_test)
```

```
In [ ]: 1
```

Need to convert the output of regression into binary outputs for calculating accuracy

```
In [111]: 1 y_pred_linear_new = [0 if x <0.5 else 1 for x in y_pred_linear]
```

```
In [112]: 1 print(confusion_matrix(y_test, y_pred_linear_new))
```

```
[[97  0]
 [ 6 89]]
```

Now we do the same process for logistic regression

```
In [113]: 1 clf = LogisticRegression()
2 scores = cross_val_score(clf, x_train, y_train, cv=5)
3 # scores = cross_val_score(clf, X, y, cv=5)
4 print("average cross validation accuracy is ",scores.mean())
5 clf.fit(x_train, y_train)
6 y_pred_logistic = clf.predict(x_test)
7 y_pred_logistic_new = [0 if x < 0.5 else 1 for x in y_pred_logistic]
8 print(confusion_matrix(y_test, y_pred_logistic_new))

average cross validation accuracy is 0.9773442226255293
[[97  0]
 [ 6 89]]
```

After doing the corss validation, average accuracy for logistic regression is at 97% which is more than linear regression's 90%.

Confusion matrix can be calculated for each model so we can calculate it for each fold in the corss validation. However here the confusion matrix was calculated directly using the testing data without involving the cross validation. Due to this, there is a difference in the cross validation accuracy reported on the training data (90% vs 97%) while the confusion matrix over the testing data remains the same. The testing data was not involved in the CV process. The confusion matrix has 192 elements which is 25% of the total 768 elements as we used standard 80-20 split for the train and testing data.